

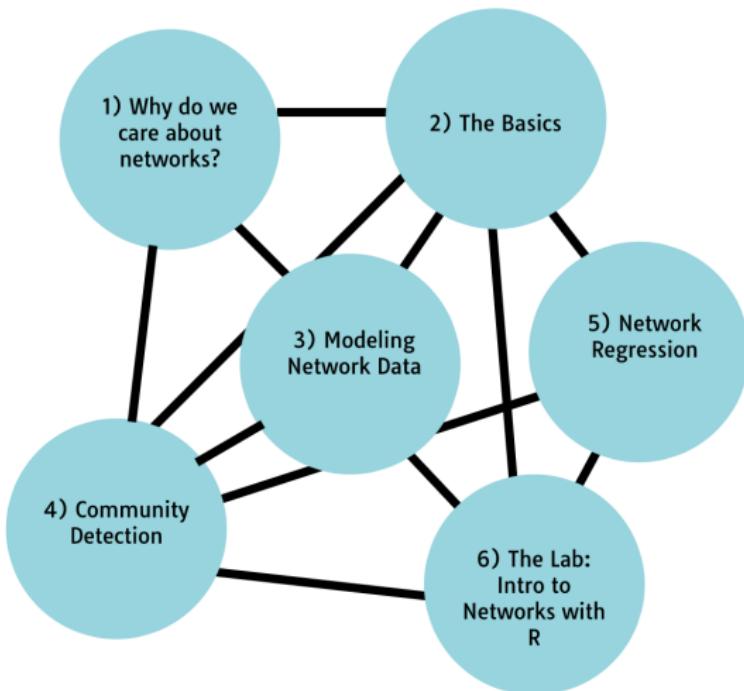
Introduction to Network Analysis

Heather Mathews

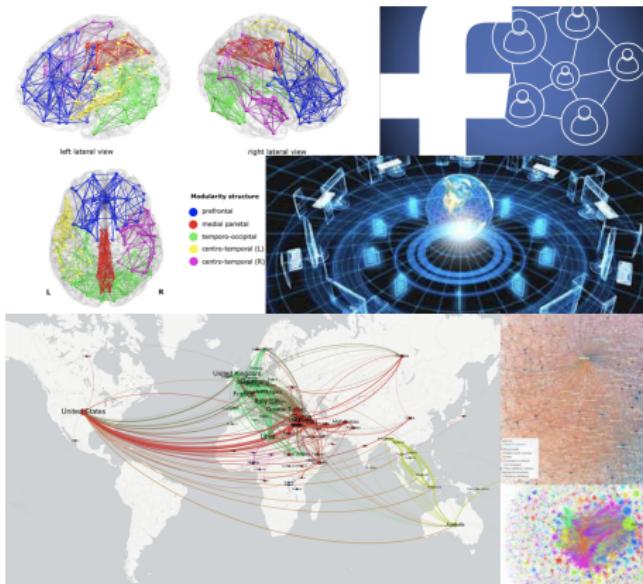
Duke University

November 7, 2019

Outline



Who uses networks?




- Almost everyone! From economists to sociologists to biologists
- Brain networks, social networks, computer networks, traffic networks, trade networks...

Motivating Example: National Longitudinal Study of Adolescent to Adult Health (AddHealth)

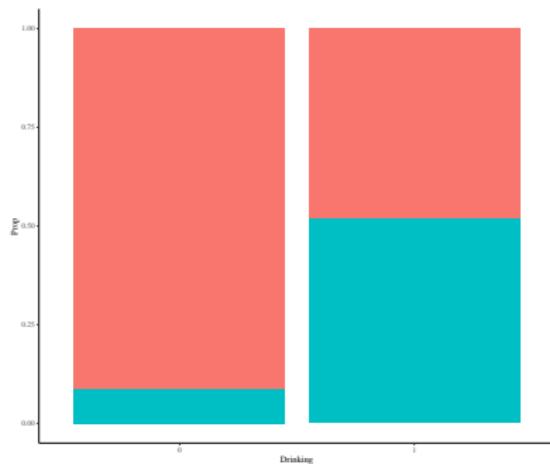
- Study was conducted due to a Congressional mandate to study factors influencing health behaviors of adolescents
- Collected data on approximately 20,000 high school students across the United States during the 1994-1995 school year
- Network Data: In each high school, each student was asked to nominate their top 5 male and female friends
- Covariate information on students was collected including grade, smoking status, drinking habits, club involvement, GPA, and sport involvement

AddHealth: Visualizing our Network

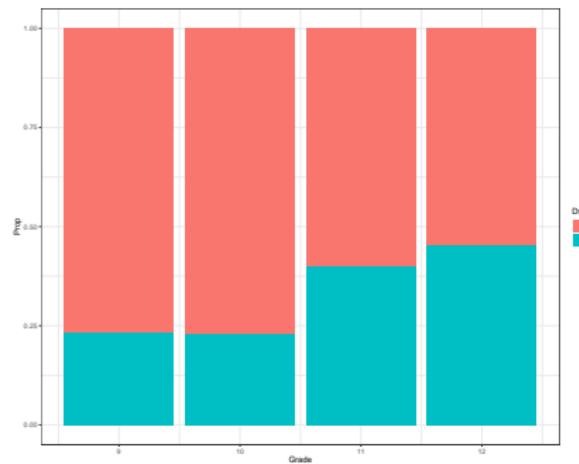


AddHealth

- Fairly uniform grade distribution
- Approximately 84% of male students are white
- About 22% of students are smokers, 30% drink alcohol
- On average, 4 individuals live in a household



Smoking and drinking



Drinking by grade

Goals of AddHealth

- Investigate which covariates might influence the probability of a friendship
- Identify possible clustering within our network. Do latent communities exist?
- Connect covariate estimation and latent communities
 - ▶ Higher GPA might increase popularity within the 'nerdy' clique but decrease popularity within the 'popular' clique



and I got that
red lip
CLASSIC

Karate Club Data

Network Scientists with Karate Trophies



5 MONTHS AGO
#NETWORKSCIENCE
#KARATECLUB
#TROPHY




The first scientist at any conference on networks who uses Zachary's karate club as an example is inducted into the Zachary Karate Club Club, and awarded a prize. This tumblr records those moments.

RSS
 ARCHIVE

Motivating Example: Karate Club

- Over three years (1970-1972), Zachary studied 34 individuals who once belonged to one karate club
- During that time, there was a conflict between John A. (administrator) and Mr. Hi (instructor) that resulted in the one club splitting into 2
- Interactions outside of the karate clubs were observed by Zachary
- He was able to classify 33/34 of the members into either John's or Mr. Hi's club
- Goal: Investigating possible fission in small community setting and interested in how information flows between the 2 clubs
- This gives us ground truth for communities! Which has made it very popular for testing community detection algorithms

Motivating Example: Karate Club



Motivating Example: Sports Analytics




What is a network?

- A network consists of **relational data** (data that describes relationships between actors)
- **Sociomatrix:** $n \times n$ matrix, A , to represent relationships between nodes. If binary entries, this is called an **adjacency matrix**

$$\begin{pmatrix} 0 & 1 & 0 & 0 & 1 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 1 & 1 \\ 1 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 \end{pmatrix}$$

Adjacency matrix, A



Graph of A

Formally Defining a Network

- **Graph:** $G = (V, E)$ with $V = \{1, \dots, n\}$ and $E = \{(i, j) \mid 1 \leq j \neq i \leq n\}$. $E \subset \mathcal{E}$ (if $E = \mathcal{E}$, fully connected graph)
 - ▶ Number of nodes: $n = |V|$
 - ▶ Number of edges: $m = |E|$

From	To
1	2
1	3
2	1
2	3
3	1

Example Edgelist

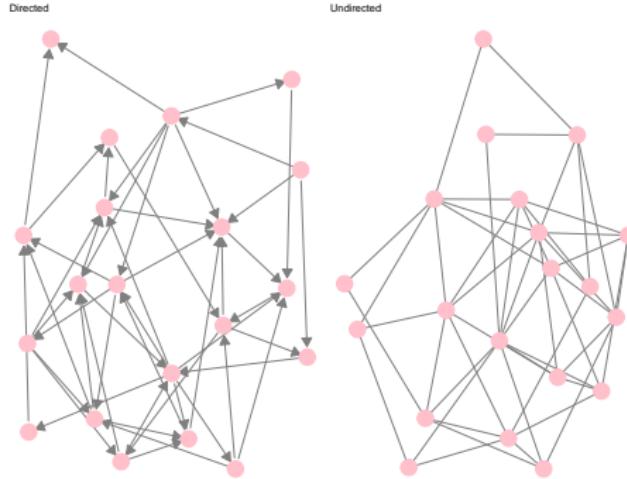
Directed vs Undirected

- Sometime edges between individuals are reciprocated (that is, if i is friends with j , j is friends with i)
- If all edges in our graph are reciprocated, then we have an **undirected network**



- However, if i is connected to j and j does not reciprocate that connection, then we have a directed edge from i to j . This leads to a **directed network**

Directed vs. Undirected



Directed: Asymmetric adjacency matrix, Undirected: Symmetric adjacency matrix

Directed vs. Undirected

Degree: Number of nodes a person is connected to

Density: Proportion of edges in graph over maximum possible number of edges

	Directed	Undirected
Max Possible # of Edges	$n^2 - n = n(n - 1)$	$n(n - 1)/2$
Degree	$d_i^{out} = \sum_{j:i \neq j} A_{i,j}$ (out) $d_i^{in} = \sum_{j:i \neq j} A_{j,i}$ (in)	$d_i = \sum_{j:i \neq j} A_{i,j}$
Density	$m/(n(n - 1))$	$2m/(n(n - 1))$

$$\begin{pmatrix} 0 & 1 & 0 & 0 & 1 \\ 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ 1 & 0 & 0 & 1 & 0 \end{pmatrix}$$

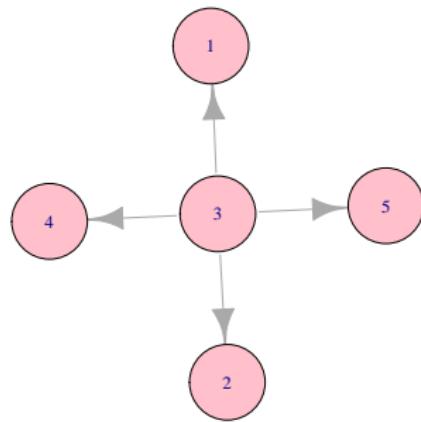
Directed matrix, A

$$\begin{pmatrix} 0 & 1 & 0 & 0 & 1 \\ 1 & 0 & 1 & 0 & 0 \\ 0 & 1 & 0 & 1 & 0 \\ 0 & 0 & 1 & 0 & 1 \\ 1 & 0 & 0 & 1 & 0 \end{pmatrix}$$

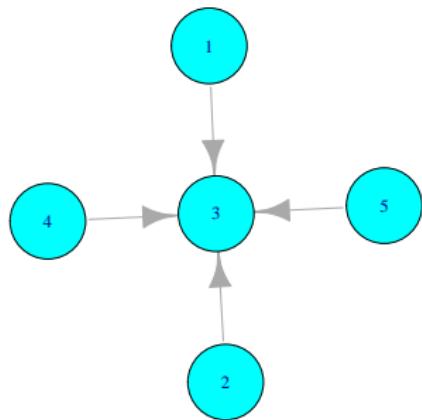
Undirected matrix, A

Sociability vs Popularity

Node 3 is Social



Node 3 is Popular



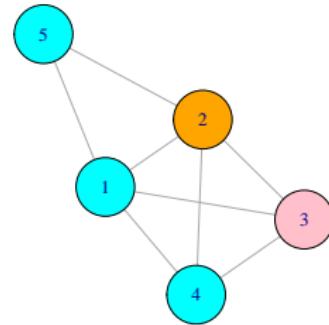
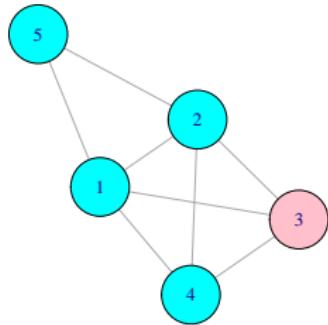
Some Relevant Definitions

- **Reciprocity:** If person 1 is connected to person 2, person 2 is connected to person 1



- **Homophily:** More likely to connect with people who are similar to you
- **Transitivity:** Friends of friends have higher probability of being friends

Sorry to say, but your friends have more friends than you...



	1	2	3	4	5
Degree	4.00	4.00	3.00	3.00	2.00
Avg. Deg. of Friends	3.00	3.00	3.67	3.67	4.00

Notions of Centrality

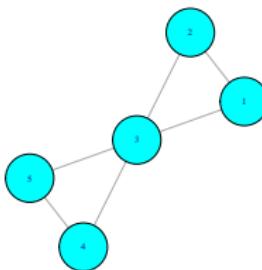
- **Betweenness Centrality:** Looks at how many times a node is part of the shortest path between other nodes

$g_{j,k}$ = # of shortest paths to get from j to k

$g_{j,k}(i)$ = # of shortest paths from j to k that go through i

$$c_i = \sum_{j < k} g_{j,k}(i) / g_{j,k}$$

- ▶ This is useful for finding individuals that are like bridges (flow of info)



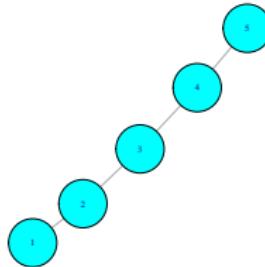
Betweenness Centrality: 0 0 4 0 0

Notions of Centrality

- **Closeness Centrality:** Measures how close one node is to all other nodes in the network. Define $d_{i,j}$ as the minimum path length from i to j .

$$c_i = \frac{1}{\sum_{j:j \neq i} d_{i,j}}$$

- ▶ Sum up the shortest paths between all nodes (good for looking at who influences spread of info)



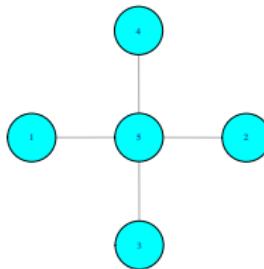
Closeness Centrality: 0.40 0.57 0.67 0.57 0.40

Notions of Centrality

- **Degree Centrality:** Importance based on number of connections a node has

$$c_i = \sum_{j:j \neq i} A_{i,j}$$

- ▶ Useful for revealing direct connections and locating popular nodes



Degree Centrality 1 1 1 1 4

Notions of Centrality


- **EigenCentrality:** Centrality of each node is proportional to the sum of its neighbor's centralities

$$c_i = \frac{1}{\lambda} \sum_{j:j \neq i} A_{i,j} c_j$$


- ▶ λ corresponds to the greatest eigenvalue of A and c corresponds to the top eigenvector. The i^{th} component of c gives the relative centrality score of vertex i
- ▶ Central nodes are connected to other central nodes (very similar to degree centrality)
- ▶ Basis for Google's PageRank
- ▶ For graph on previous slide, eigen centrality is: 0.5 0.5 0.5 0.5 1.0

Example: Karate Club


Degree




Closeness



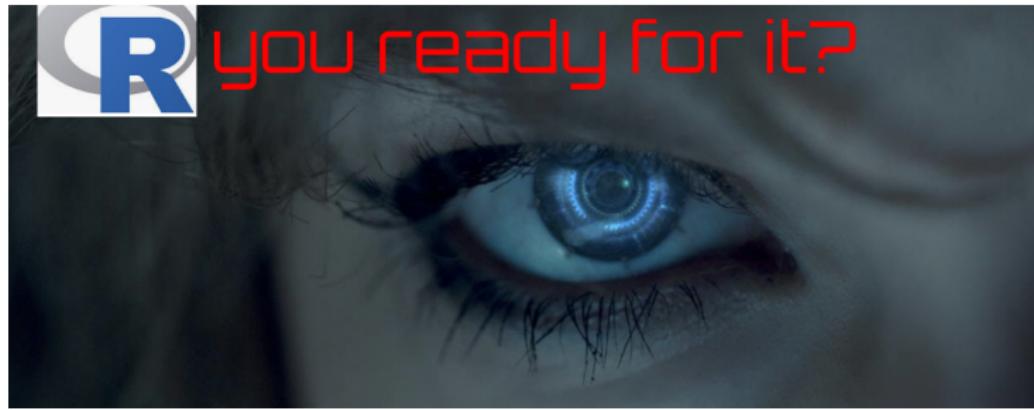
Eigen



Betweenness



Now off to the lab...



Modeling and Generating Networks

- Erdos-Renyi Graph (ER)
- Exponential Random Graph Models (ERGMS)
- Stochastic Block Models (SBM)
- Additive and Multiplicative Effects Network (AMEN)

Erdos-Renyi Graph

Generating graphs:

- **G(n,p)**: ER graph where edges are drawn independently with probability, p

$$P(G_0) = P(G = G_0) = p^m(1 - p)^{(N-m)}$$

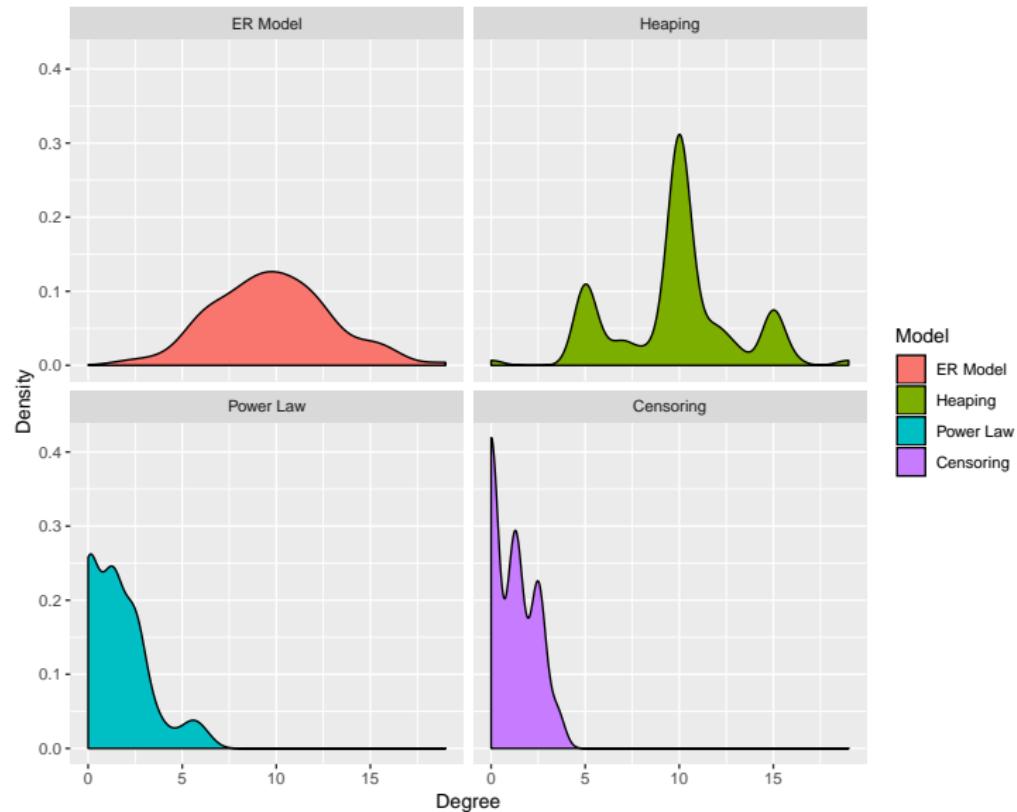
$$E(\# \text{ of Edges}) = \binom{n}{2}p$$

Relating ER to the Adj. Matrix

How do the models on the previous slide relate to the adjacency matrix, A ?

$$A_{i,j}|p \sim Bern(p)$$

Degree Distribution... and why we care



Degree Distribution

- What is the degree distribution of an ER model?
- ER(p) has $(n-1)$ possible friends
- Let person i have k friends, thus $\binom{n-1}{k}$ possibilities and probability of a friend is p :

$$\binom{n-1}{k} p^k (1-p)^{n-1-k}$$

- Bin($n-1, p$)

Back to the lab...



Exponential Random Graph Models: ERGM

Basic generative model for networks that is based off of sufficient statistics

$$P_{\theta}(G_n) = \exp(\theta^T t_n(G_n) - \Psi_n(\theta))$$

where

- θ : Parameter (or vector of parameters) we want to estimate
- t_n : Sufficient statistic
- Ψ_n : Normalizing constant

Examples of ERGM

The ER Model:

- In this model, each edge is sampled iid Bernoulli with some probability p . For this model, $t_n = \sum_{i < j} A_{i,j}$

$$P_\theta(A_n) \propto \exp\left\{\theta \sum_{i < j} A_{ij}\right\}$$

ER as an ERGM

We have seen the likelihood for the $ER(p)$:

$$\begin{aligned} p^m(1-p)^{N-m} &= \exp\{\log(p^m(1-p)^{N-m})\} \\ &= \exp\{m \times \log(p) + (N-m) \times \log(1-p)\} \\ &= \exp\left\{m \times \log\left(\frac{p}{1-p}\right) + N \times \log(1-p)\right\} \end{aligned}$$

Recall the ERGM

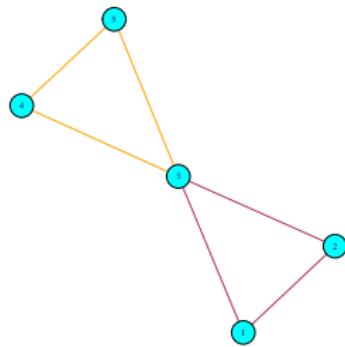
$$P_\theta(A_n) \propto \exp(\theta t_n(A_n)) \quad (1)$$

where


- $\theta = \log\left(\frac{p}{1-p}\right)$
- $t_n = m = \sum_{i < j} A_{i,j}$

ERGM Cont.

We can also consider some other sufficient statistics to include in our model such as # of stars, # of triangles, # of edges, etc.



Triangles



Star

What kinds of questions can ERGMs help answer?

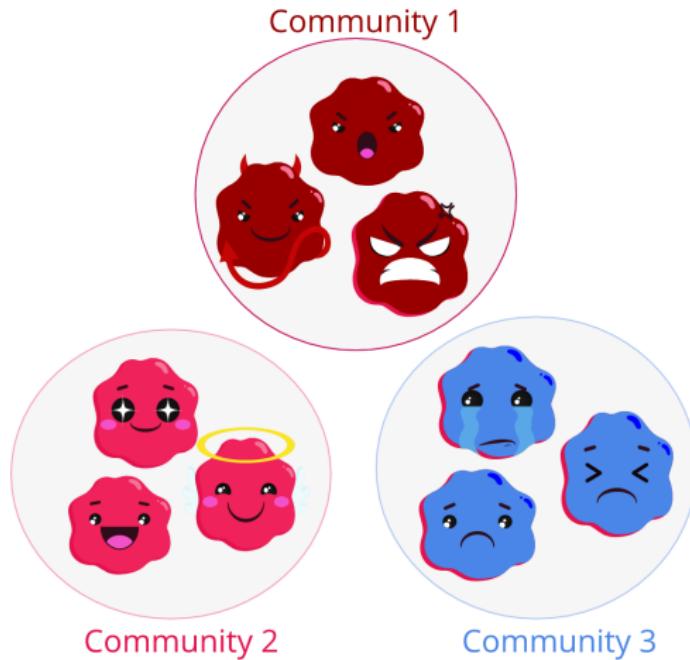
- We can look at expectation of degree, edges, degree distribution
- We can test what model fits a new network best (were 2 graphs generated from the same model?)
- HOWEVER, not always consistent estimators :/

Back to the lab...



The Stochastic Block Model


Now we consider if nodes come from **communities** (Holland et al, 1983)



$$\theta = \begin{pmatrix} 1/3 & 1/3 & 1/3 \end{pmatrix}$$

Understanding the SBM Visually

$$B = \begin{pmatrix} 0.4 & 0.05 & 0.3 \\ 0.05 & 0.4 & 0.1 \\ 0.3 & 0.1 & 0.4 \end{pmatrix}$$



$$P(\text{Angel} \rightarrow \text{Devil} | \text{Angel}, \text{Devil}) = 0.05$$

Another Generative Model: The Stochastic Block Model

- $\theta = [\theta_1, \theta_2, \dots, \theta_K]$ is a vector containing probabilities that a person belongs to a particular community, $k \in \{1, \dots, K\}$
- $B \in \mathbb{R}^{K \times K}$ is a preference matrix that describes the probability of connection of nodes based solely on an individual's membership
- Z indicates which community a person belongs to

$$\begin{aligned} P(Z_i = k) &= \theta_k \\ A_{i,j} | Z_i, Z_j &\sim \text{Bern}(B_{Z_i, Z_j}) \\ P(A, Z, \theta, B) &= \prod_K \theta_k^{\sum 1_{z_i=k}} \prod_{i,j} B_{Z_i, Z_j}^{A_{ij}} (1 - B_{Z_i, Z_j})^{1-A_{ij}} \end{aligned}$$

Balanced Multi-Label Propagation for Overlapping Community Detection in Social Networks

Authors

Authors and affiliations

Zhi-Hao Wu , You-Fang Lin, Steve Gregory, Huai-Yu Wan, Sheng-Feng Tian

Mutually Enhancing Community Detection and Sentiment Analysis on Twitter Networks

William Deitrick, Wei Hu*

Bayesian Inference and Testing of Group Differences in Brain Networks

Listen

[Listen to this page using the media player](#)

Identifying functional urban regions within traffic flow
Ed Markey

Pages 40-42 | Received 30 Jan 2014, Accepted 03 Feb 2014, Published online: 12 Mar 2014

Daniele Durante* and David B. Dunson†

Promoting Small and Medium Enterprises with a Clustering Approach: A Policy Experience from Indonesia by Tulus Tambunan

Community Detection in General Stochastic Block models: Fundamental Limits and Efficient Algorithms for Recovery

Publisher: IEEE

2 Author(s)

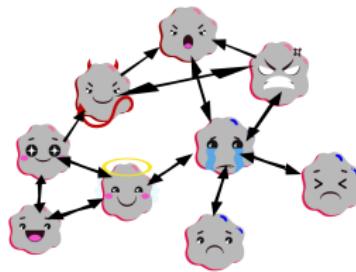
Emmanuel Abbe ; Colin Sandon

[View All Authors](#)

Social selection and peer influence in an online social network

Community Detection

- Maybe we believe our network, A , came from an SBM, and we care about finding community labels



Observed Network



Methods for Community Detection

- Spectral methods
- Gaussian Mixture Models
- Centrality based approaches

Spectral Clustering


- Consider taking a *spectral decomposition* of A (eigendecomposition, singular value decomposition)
- If we take the eigendecomposition,

$$A = V \Lambda V^T$$


where $\Lambda = \text{Diag}(\lambda_1, \dots, \lambda_n)$ contains the eigenvalues in decreasing order (in magnitude), and $V = (V_1, \dots, V_n)^T$ is a matrix containing the orthonormal columns corresponding to the eigenvectors

- If community structure exists, it should presumably show up in a lower dimensional representation of A
- Consider taking the top K eigenvectors of A where K is the number of hypothesized communities that exist

Looking at our eigenvectors



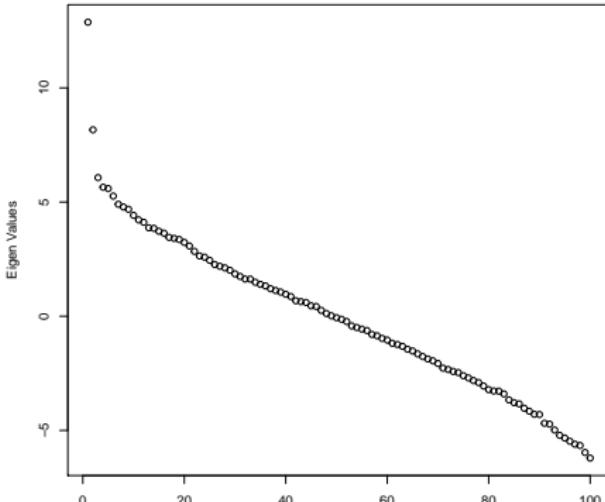
True Comm.
● 1
● 2



Est K-means Comm.
● 1
● 2

The Ultimate Question: How do we pick the number of communities?



- The answer to this is not super clear
- Elbow plots, look for eigen values that escape the bulk
- Prior information
- Try a few different options



Once we have our number of communities...

[Rohe et al., 2012]

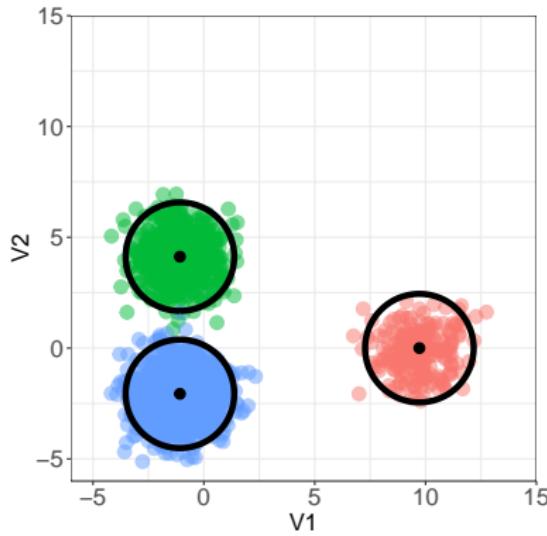
- Run a clustering algorithm on the top K eigenvectors corresponding to the top K eigenvalues (in magnitude), $V_{:,1:K}$
- Can also cluster on $V_{:,1:K}\Lambda_{1:K,1:K}^{1/2}$



Gaussian Mixture Models

- Consider the latent positions that we can obtain from spectral decompositions of A
- It can actually be shown that these latent positions can be modeled using Gaussian Mixture Models:

$$P(V_i = v_i) = \prod_{i=1}^n \sum_{k=1}^K \pi_k N(v_i; \mu_k, \sigma_k^2)$$




Moving to Centrality Based Example...




Edge Betweenness Algorithm

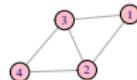
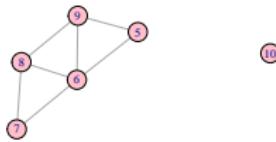
- Basic Idea:
 - ▶ Find the edge with maximum edge betweenness centrality and delete it



Step 1



Edge Betweenness Algorithm



Step 4



Back to the lab...



Moving to Regression: Connecting to things we know

- Consider a typical probit regression where we observe a binary response and covariates
- We want to estimate β so we consider probit regression where ϵ are assumed to be iid

$$A = \mathbb{1}_{Z>0}$$

$$Z = X\beta + \epsilon$$

Connecting to things we know

A is a type of response variable, but it is a matrix. To get something more familiar, we could vectorize A such that:

$$A = \begin{pmatrix} 0 & 1 & 0 & 1 & 1 \\ 1 & 0 & 1 & 1 & 0 \\ 1 & 0 & 0 & 0 & 1 \\ 0 & 1 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 \end{pmatrix} \rightarrow \text{vec}(A) = \begin{pmatrix} A_{:,1} \\ A_{:,2} \\ A_{:,3} \\ A_{:,4} \end{pmatrix} = \begin{pmatrix} 0 \\ 1 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ \vdots \\ 0 \\ 0 \end{pmatrix}$$

Remember that dyadic dependence...



However, with a network, it is not reasonable to assume that all of these entries are independent

Social Relations Model [Warner et al., 1979]

- We are interested in modeling the variability in A

$$a_{i,j} = \mathbb{1}_{z_{i,j}>0}$$
$$z_{i,j} = \beta_0 + c_i + d_j + \epsilon_{i,j}$$

- β_0 : Overall global mean
- c : Individual row (sociability/sender behavior) random effects
- d : Individual column (popularity/reciever behavior) random effects

$$\begin{pmatrix} c_i \\ d_i \end{pmatrix} \stackrel{i.i.d}{\sim} N(0, \Sigma_{cd}) \text{ where } \Sigma_{cd} = \begin{bmatrix} \sigma_c^2 & \sigma_{cd} \\ \sigma_{cd} & \sigma_d^2 \end{bmatrix}$$

$$\begin{pmatrix} \epsilon_{i,j} \\ \epsilon_{j,i} \end{pmatrix} \stackrel{i.i.d}{\sim} N(0, \Sigma_e) \text{ where } \Sigma_e = \begin{bmatrix} 1 & \rho \\ \rho & 1 \end{bmatrix}$$

Adding in covariate information

$$a_{i,j} = \mathbb{1}_{z_{i,j} > 0}$$
$$z_{i,j} =$$
$$\beta_0 + \sum_{p=1}^{P_r} (\textcolor{blue}{x_{r,p,i}} \beta_{r,p}) + \sum_{p=1}^{P_c} (\textcolor{blue}{x_{c,p,j}} \beta_{c,p}) + \sum_{p=1}^{P_d} (\textcolor{blue}{x_{d,p,i,j}} \beta_{d,p}) + c_i + d_j + \epsilon_{i,j}$$

- $\textcolor{blue}{X_r}$: Observed covariate information for row covariates
- $\textcolor{blue}{X_c}$: Observed covariate information for column covariates
- $\textcolor{blue}{X_d}$: Observed covariate information for dyadic covariates
- β : Coefficients of interest estimating covariate effects on connections

Putting this into matrix form

$$\begin{bmatrix} z_{1,1} & z_{1,2} & z_{1,3} \\ z_{2,1} & z_{2,2} & z_{2,3} \\ z_{3,1} & z_{3,2} & z_{3,3} \end{bmatrix} = \begin{bmatrix} x_{r,1} & x_{r,1} & x_{r,1} \\ x_{r,2} & x_{r,2} & x_{r,2} \\ x_{r,3} & x_{r,3} & x_{r,3} \end{bmatrix} \beta_r + \begin{bmatrix} x_{c,1} & x_{c,2} & x_{c,3} \\ x_{c,1} & x_{c,2} & x_{c,3} \\ x_{c,1} & x_{c,2} & x_{c,3} \end{bmatrix} \beta_c + \dots$$

Better Models...



Adding in multiplicative effects (AMEN [Hoff, 2018])

- Sometimes there are higher order latent dependencies between nodes
 - ▶ **Example:** Nodes may be *homophilous*. Individuals who are similar to one another are more likely to be friends (clustering)


$$a_{i,j} = \mathbb{1}_{z_{i,j} > 0}$$

$$z_{i,j} = \beta_0 + \sum_{p=1}^{P_r} (\textcolor{blue}{x_{r,p,i}} \beta_{r,p}) + \sum_{p=1}^{P_c} (\textcolor{blue}{x_{c,p,j}} \beta_{c,p}) + \sum_{p=1}^{P_d} (\textcolor{blue}{x_{d,p,i,j}} \beta_{d,p}) + c_i + d_j + \textcolor{blue}{u_i^T v_j} + \epsilon_{i,j}$$

- $\textcolor{blue}{U}, \textcolor{teal}{V}$: Latent factor matrices of rank R ($n \times R$ matrices)
- $\textcolor{blue}{u_i}$ gives us information about a node as a sender
- $\textcolor{teal}{v_j}$ gives us information about a node as a receiver
- They can describe notions of stochastic equivalence (if $\textcolor{blue}{u_i}$ is similar to $\textcolor{blue}{u_j}$, then they may share similar behaviors)

But what if we have latent community structure?

- Latent multiplicative effects can capture latent community structure



How do we use this model?

- The standard AMEN model is implemented using a Markov Chain Monte Carlo (MCMC) algorithm
- For the standard model, this is provided in the ‘amen’ R package

BRACE YOURSELVES



memegenerator.net

Brief overview of Bayesian Methods

- In frequentist methods, we have a parameter that we want to estimate, θ , that is considered to be fixed, but unknown
- In Bayesian methods, we are still interested in estimating θ , however we believe it is an unknown, **random** quantity
- Rather than come up with a point estimate, we want to derive a posterior distribution for θ . What does that mean exactly?

Bayesian Methods

- We observe a dataset, y , which comes from a sample space, \mathcal{Y} where \mathcal{Y} represents all possible datasets that y could come from
- We are interested in estimating a population parameter, θ , that comes from Θ (parameter space of θ)
- In a Bayesian setting, we put a *prior* distribution on our parameter, $\theta \in \Theta$. This describes our beliefs about the true population parameter ($p(\theta)$)
- We then have a *sampling model*, $p(y|\theta)$ that describes our beliefs about y had we known the true θ
- Our goal is then to get $p(\theta|y)$ which describes our beliefs for the true value of θ after observing our data y

Bayes Rule

$$p(\theta|y) = \frac{p(y|\theta)p(\theta)}{\int_{\tilde{\theta} \in \Theta} p(y|\tilde{\theta})p(\tilde{\theta})d\tilde{\theta}}$$

Back to AMEN: How do we implement this model?

- We have a model: $Z \sim P(Z|\theta)$, $\theta \in \Theta$
- When Y is binary,

$$S(A) = \{Z : a_{i,j} > 0 \Rightarrow z_{i,j} > 0, a_{i,j} = 0 \Rightarrow z_{i,j} \leq 0\}$$

- Likelihood is:

$$L_B(\theta, Z) = \Pr(Z \in S(A)|\theta) = \int_{S(A)} P(Z|\theta) d\mu(Z)$$

- We can then use a Gibbs Sampler with MH to approximate $P(\theta|Z \in S(A))$

Why AMEN? Different likelihoods!

- Up to this point, we have focused on adjacency matrices/ binary networks
- However, networks can have many different forms
- AMEN allows us to easily consider some of these
 - ▶ **Censored Binary:** Still 1 or 0, but not all connections are observed
 - ▶ **Ranked:** Connections are ranked by some level of importance
 - ▶ **Fixed Rank Nomination:** A node can only have a FIXED number of connections and they are ranked

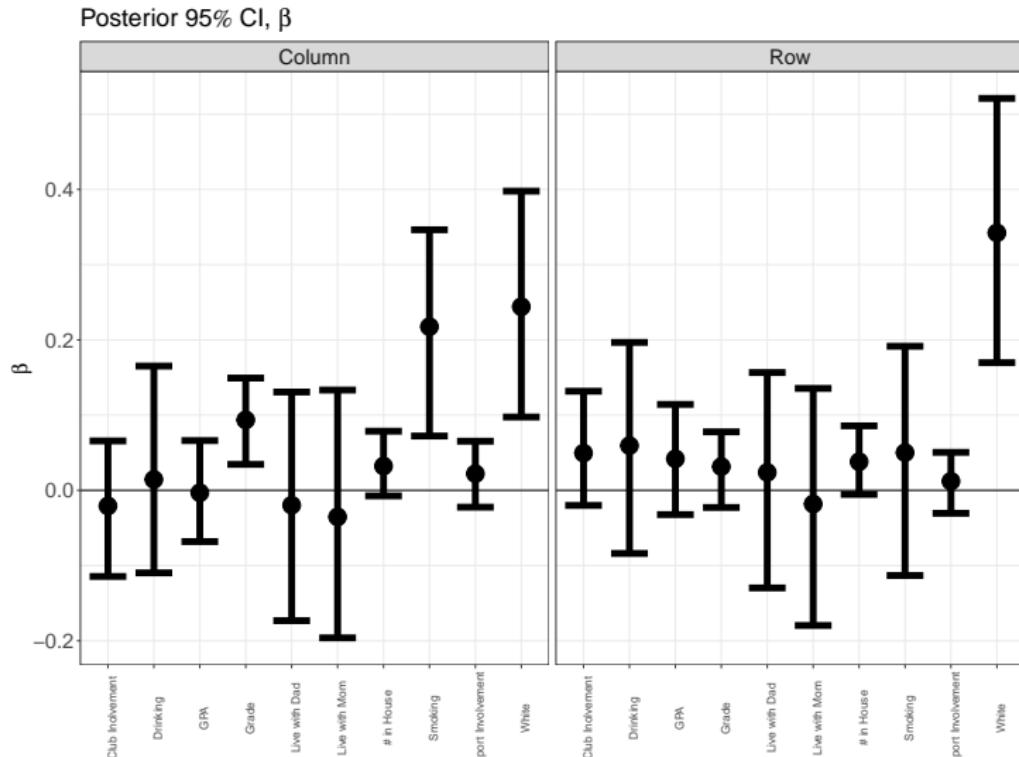
Fixed Rank Nomination

AddHealth


- Each person could rank up to 5 male friends and 5 female friends
- This introduces a *censorship* issue.
 - If person i ranks the 5 people, are they friends with the 6th and just didn't have room to rank them? or do they not like that person?
- If someone ranks less than 5 people, then we assume that they are not friends with person 6
- Another issue... perhaps person i just has never met person 6 but they would be great friends if they had

Back to AddHealth: Regression Results

- 95% CI for β estimates when fitting AMEN model with $R = 2$ on AddHealth network



Standard AMEN R=2



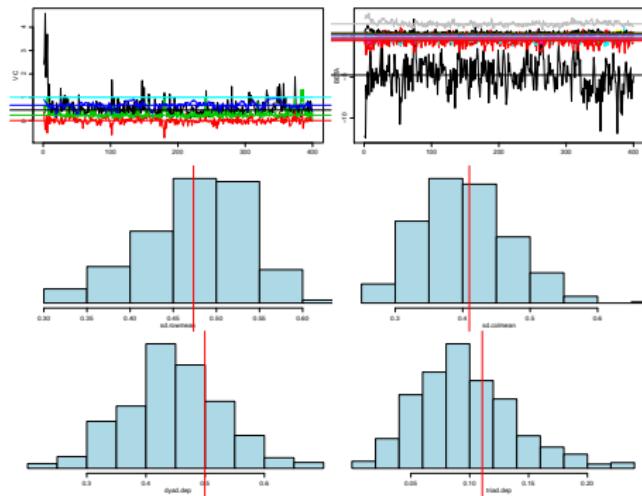
What appears to be important?

- Column: Grade, Smoking, White
- Row: White
- Dyadic: Sharing same drinking behavior, grade, smoking behavior

Diagnostics: Goodness of Fit Statistics

AMEN provides us with some posterior predictive goodness of fit statistics:

- ① Empirical standard deviation for row means
- ② Empirical standard deviation for column means
- ③ Empirical within-dyad correlation
- ④ Normalized measure of triadic dependence



And back to the lab...



<https://igraph.org/r/doc/igraph.pdf> (igraph)

<https://cran.r-project.org/web/packages/ergm/ergm.pdf> (ergm)

<https://arxiv.org/pdf/1506.08237.pdf> (amen)



Abbe, E. and Sandon, C. (2015).

Community detection in general stochastic block models: Fundamental limits and efficient algorithms for recovery.

In *2015 IEEE 56th Annual Symposium on Foundations of Computer Science*, pages 670–688. IEEE.



Binkiewicz, N., Vogelstein, J. T., and Rohe, K. (2017).

Covariate-assisted spectral clustering.

Biometrika, 104(2):361–377.



Deitrick, W. and Hu, W. (2013).

Mutually enhancing community detection and sentiment analysis on twitter networks.

Journal of Data Analysis and Information Processing, 1(03):19.



Durante, D., Dunson, D. B., et al. (2018).

Bayesian inference and testing of group differences in brain networks.

Bayesian Analysis, 13(1):29–58.



Feld, S. L. (1991).

Why your friends have more friends than you do.

American Journal of Sociology, 96(6):1464–1477.



Harris, K., Halpern, E., Whitsel, E., Hussey, J., and Udry, J. (2009).

The national longitudinal study of adolescent health: Research design.



Hoff, P. (2018).

Additive and multiplicative effects network models.

arXiv preprint arXiv:1807.08038.



Hoff, P., Fosdick, B., Volfovsky, A., and He, Y. (2017).

amen: Additive and Multiplicative Effects Models for Networks and Relational Data.

R package version 1.3.



Hoff, P., Fosdick, B., Volfovsky, A., and Stovel, K. (2013).

Likelihoods for fixed rank nomination networks.

Network Science, 1(03):253–277.



Hoff, P., Raftery, A., and Handcock, M. (2002).

Latent space approaches to social network analysis.

Journal of the american Statistical association, 97(460):1090–1098.



Holland, P., Laskey, K., and Leinhardt, S. (1983a).

Stochastic blockmodels : First steps.

Social Networks, 5:109–137.



Holland, P., Laskey, K., and Leinhardt, S. (1983b).

Stochastic blockmodels: First steps.

Social Networks, 5:109–137.



Holland, P. and Leinhardt, S. (1981).

An exponential family of probability distributions for directed graphs.

Journal of the American Statistical Association, 76:33–50.



Lewis, K., Gonzalez, M., and Kaufman, J. (2012).

Social selection and peer influence in an online social network.

Proceedings of the National Academy of Sciences, 109(1):68–72.

-  Manley, E. (2014).
Identifying functional urban regions within traffic flow.
Regional Studies, Regional Science, 1(1):40–42.
-  Mossel, E., Neeman, J., and Sly, A. (2014).
Belief propagation, robust reconstruction and optimal recovery of block models.
In *Conference on Learning Theory*, pages 356–370.
-  Rohe, K., Chatterjee, S., and Yu, B. (2011).
Spectral clustering and the high -dimensional stochastic blockmodel.
The Annals of Statistics, 39(4):1878–1915.
-  Rohe, K., Qin, T., and Yu, B. (2012).
Co-clustering for directed graphs: the stochastic co-blockmodel and spectral algorithm di-sim.
arXiv preprint arXiv:1204.2296.

-  Shiokawa, H., Fujiwara, Y., and Onizuka, M. (2013).
Fast algorithm for modularity-based graph clustering.
In *Twenty-Seventh AAAI Conference on Artificial Intelligence*.
-  Tambunan, T. (2005).
Promoting small and medium enterprises with a clustering approach: A policy experience from indonesia.
Journal of Small Business Management, 43(2):138–154.
-  Warner, R. M., Kenny, D. A., and Stoto, M. (1979).
A new round robin analysis of variance for social interaction data.
Journal of Personality and Social Psychology, 37(10):1742.
-  Wu, Z.-H., Lin, Y.-F., Gregory, S., Wan, H.-Y., and Tian, S.-F. (2012).
Balanced multi-label propagation for overlapping community detection in social networks.
Journal of Computer Science and Technology, 27(3):468–479.